

The Temporal Tug-of-War: Visualizing and Detecting RAG Conflicts in Diffusion Models via Trajectory Variance

Anonymous ACL submission

Abstract

Retrieval-Augmented Generation (RAG) introduces a failure mode in Discrete Diffusion Language Models (DLMs): when a retrieved context contradicts the model’s parametric memory, the iterative denoising process becomes an observable battleground between competing knowledge sources. We propose the **Trajectory Variance Score (TVS)**, a simple and interpretable method for detecting these RAG conflicts. By computing the mean pairwise cosine distance of answer embeddings across independent stochastic denoising trajectories, TVS captures the temporal “tug-of-war” between parametric and contextual attractors. Our method is computationally lightweight, requiring only $N = 2$ parallel inference runs. Furthermore, we demonstrate that a basic Logistic Regression classifier achieves competitive conflict detection accuracy without requiring complex sequential modeling. We evaluate TVS across four diverse datasets (Synthetic, SciQ, PopQA, CounterFact) on two architecturally distinct DLMs—LLaDA and Dream 7B—demonstrating that the TVS signal and its key properties transfer across diffusion architectures.

1 Introduction

The adoption of Retrieval-Augmented Generation (RAG) (Lewis et al., 2020) has become the standard mechanism for mitigating hallucinations in large language models. However, this architecture introduces a secondary failure mode: knowledge conflict. When an injected RAG context directly contradicts the model’s pre-trained parametric memory (e.g., outdated facts or adversarially generated fakes), the model must implicitly arbitrate between its internal prior and the external evidence.

In standard autoregressive (AR) models, this friction is instantaneous and localized, manifesting as spikes in token-level perplexity during left-to-right decoding. However, the emergence of Discrete Dif-

fusion Language Models (DLMs) like LLaDA (Nie et al., 2025) fundamentally alters this paradigm. Unlike AR models, DLMs generate text through an iterative, non-causal denoising process. The resolution of knowledge conflicts is therefore not a single token-emission event, but a temporal “tug-of-war” distributed across the entire diffusion trajectory.

Recent work in DLM interpretability, such as TraceDet (Chang et al., 2025), has demonstrated that features extracted from hidden trajectories can classify intrinsic hallucinations. Building on this view, we introduce the **Trajectory Variance Score (TVS)**. We hypothesize that forced arbitration between contradictory memory sources induces measurable instability in the intermediate denoising states—and that this instability is directly observable as elevated cross-seed semantic divergence. By executing N independent stochastic denoising runs and computing mean pairwise cosine distance across seed predictions at each timestep, we obtain a temporal feature vector that is both interpretable and classifiable by simple models.

Our contributions are:

1. We formally define the problem of RAG-induced knowledge friction in the discrete diffusion paradigm.
2. We propose **TVS**, computed as the mean pairwise cosine distance of sentence-embedded predictions across N seeds at each denoising timestep, providing an interpretable, logit-free measure of cross-seed semantic divergence over the stochastic trajectory.
3. We demonstrate that the TVS signal does not require complex sequential modeling: a simple Logistic Regression classifier over the flattened T -dimensional TVS vector achieves competitive conflict detection accuracy, matching the performance of recurrent architectures.

4. We validate the generality of TVS across two architecturally distinct DLMs—LLaDA (low-confidence remasking) and Dream 7B (entropy-based remasking)—showing that the three-phase temporal dynamics, the sufficiency of linear classifiers, and the cross-dataset difficulty ordering all transfer, while revealing architecture-dependent seed sensitivity.

2 Related Work

2.1 Hallucination and Conflict Detection in AR-LLMs

Hallucination detection in autoregressive models has been extensively studied. Output-based detection leverages signals derived from the generated text, such as semantic entropy (Kuhn et al., 2023) or lexical similarity. The core intuition is that hallucinated or conflicting responses are associated with lower predictive confidence. Latent-based methods instead probe the hidden representations during a single forward pass, utilizing techniques like Contrast-Consistent Search (CCS) (Burns et al., 2023) to separate truthful from hallucinated states. However, these methods are uniquely suited to the left-to-right generation paradigm and fail to capture the multi-step refinement of diffusion models.

2.2 Discrete Diffusion Language Models

Discrete Diffusion Language Models have emerged as alternatives to AR models (Sahoo et al., 2024; Shi et al., 2024). Models such as LLaDA (Nie et al., 2025) utilize discrete remasking processes to scale up to 8B parameters, achieving performance comparable to leading LLMs. Similarly, DREAM (Ye et al., 2025) adapts diffusion paradigms for enhanced reasoning. Recent efforts like SPREAD (Yu et al., 2026) have begun exploring DLMs within the RAG framework, identifying Response Semantic Drift (RSD) and employing query-relevance-guided denoising to keep generation anchored to the query semantics. Furthermore, while recent work like Adaptive Retrieval-Augmented Masked Diffusion (ARAM) (Kim and Ye, 2026) attempts to force alignment via adaptive logit guidance, our work is fundamentally different because it provides a post-hoc interpretability window without altering the generation process. Consequently, the specific mechanisms of detecting the underlying parametric versus external knowledge conflicts remain underexplored.

2.3 Trajectory Analysis in Diffusion Models

The shift to iterative refinement necessitates new interpretability tools.

TraceDet (Chang et al., 2025) is the most relevant precursor to our work. It formulates the denoising process as an action trace and applies the Variational Information Bottleneck principle to extract compressed sub-traces that are maximally informative of hallucinated outputs. TraceDet operates on the model’s internal hidden states, requiring access to intermediate representations across the full trajectory. In contrast, TVS is entirely external: it requires only decoded text outputs from parallel runs, making it model-agnostic and applicable without architectural modifications.

TDGNet (Hemmat et al., 2026) constructs evolving token-level attention graphs at each denoising timestep and learns graph neural network representations over these temporal dynamic graphs to detect hallucinations. While effective, TDGNet’s reliance on fine-grained attention patterns makes it computationally expensive and tightly coupled to the model’s internal attention mechanism—a dependency that TVS avoids by operating purely in sentence embedding space.

OSCAR (Shah et al., 2026) computes cross-chain Shannon entropy across parallel denoising chains with randomized reveal orders to detect high-uncertainty token positions. While OSCAR and TVS both exploit parallel denoising runs, they differ fundamentally: OSCAR operates token-by-token within a single pass to flag high-entropy positions; TVS measures the *temporal evolution of cross-run semantic divergence* in the answer entity embedding space—a complementary signal that captures how competing knowledge sources destabilize generation dynamics over the full denoising schedule.

None of these methods address the specific phenomenon of RAG-induced knowledge friction, nor demonstrate that the resulting signal can be classified without complex sequential modeling.

3 Background: Diffusion Large Language Models

3.1 Formal Formulation of D-LLMs

Unlike AR models that rely on the next-token prediction paradigm $P(x_i|x_{<i})$, DLMs generate responses by iterative refinement through a forward noising and backward denoising process over T time steps.

Let $r = (r_0, \dots, r_T)$ denote the sequence of intermediate texts, where each $r_t \in \mathcal{V}^n$ is a token sequence of length n . Here, r_0 represents the fully clean, structured text sequence, and r_T represents the fully masked sequence (pure noise).

Given an input query p_0 , the **forward noising process** is defined by a sequence of transition distributions $\{q(r_t|r_{t-1})\}_{t=1}^T$, mathematically expressed as:

$$q(r_{1:T}|r_0) = \prod_{t=1}^T q(r_t|r_{t-1}) \quad (1)$$

This process gradually corrupts the sequence by replacing tokens with a designated ‘[MASK]’ state.

Conversely, the **reverse denoising process** is parameterized by a neural network defining a sequence of conditional distributions $\{P_\theta(r_0|r_t, p_0)\}_{t=1}^T$. At each timestep $t > 0$, the model predicts all masked tokens simultaneously from the current state r_t , yielding an estimated clean sequence $\tilde{r}_{t-1} \sim P_\theta(r_0|r_t, p_0)$. Following this prediction, a discrete remasking strategy is applied (typically low-confidence remasking) where a specific fraction ρ_t of the predicted tokens are remasked to form the intermediate state r_{t-1} .

3.2 Knowledge Friction in the Denoising Trajectory

We model the generation process as a stochastic denoising trajectory $\tau = (r_0, r_1, \dots, r_T)$, where each intermediate state r_t represents a partially-unmasked sequence. Under a clean RAG context, the trajectory converges smoothly toward a single semantic attractor. Under a conflicting context, competing knowledge sources (the parametric prior θ_{true} and the injected context p_0) create sustained instability in the intermediate states, a phenomenon we term **Knowledge Friction** (Yu et al., 2026).

For example, given the query “What is the capital of France?” the model’s prior strongly associates the answer with “Paris” (θ_{true}). If the injected RAG context states “The capital of France is London” (p_0), the stochastic denoising trajectory oscillates between these competing attractors before crystallizing. Because DLMs decode all positions in parallel, this friction is distributed globally across the full denoising schedule rather than localized to a single token emission event.

4 Methodology: Trajectory Variance Score

To detect this knowledge friction, we isolate the temporal instability it creates by measuring the semantic divergence across multiple independent stochastic denoising runs. The end-to-end pipeline is illustrated in Figure 1.

End-to-End Evaluation Pipeline for Knowledge Friction




Figure 1: End-to-end evaluation pipeline: facts undergo parametric memory filtering, conflict contexts are injected, $N = 2$ independent denoising trajectories are executed, the mean pairwise cosine distance (TVS) is computed at each of $T = 50$ timesteps, producing a T -dimensional feature vector, and the resulting vector is classified by a Logistic Regression detector.

4.1 Parametric Memory Filtering and Trajectory Generation

Diffusion models are inherently stochastic; each denoising trajectory depends on the initial random seed governing the remasking schedule. Before generating conflict trajectories, we apply a rigorous parametric memory filter: a fact is retained only if the base DLM correctly generates the true answer across at least 4 out of $N = 5$ independent seeds. This ensures subsequent friction is genuinely induced by the RAG context rather than inherent model ignorance. After filtering, **2,947 facts** survived for LLaDA evaluation (yielding 5,894 balanced clean/conflict sample pairs).

For each verified fact, we execute $N = 2\text{--}5$ independent stochastic denoising runs over $T = 50$ diffusion timesteps. At each run i and timestep t , we decode the model’s predicted token sequence. We isolate the answer entity from the full generation via exact sub-string matching against the target, and embed it using the all-mpnet-base-v2 sentence encoder to obtain embedding $\mathbf{e}_{i,t} \in \mathbb{R}^d$.

4.2 TVS Formulation

In the absence of knowledge friction, all N runs rapidly converge on identical semantic structures, yielding low cross-seed divergence. Under conflict, competing knowledge sources cause the denoising trajectories to diverge before finally crystallizing.

At each timestep t , we isolate the predicted answer entity from each seed’s decoded text, embed it using a frozen sentence encoder, and compute the mean pairwise cosine distance across the N seed embeddings:

$$\text{TVS}_t = \frac{1}{\binom{N}{2}} \sum_{i < j} (1 - \cos(\mathbf{e}_{i,t}, \mathbf{e}_{j,t})) \quad (2)$$

where $\mathbf{e}_{i,t} \in \mathbb{R}^d$ is the sentence embedding of the predicted entity at seed i , timestep t , and $\cos(\cdot, \cdot)$ denotes cosine similarity. This formulation directly measures semantic divergence in embedding space, making it interpretable and independent of the model’s internal logits.

This yields a feature vector $\mathbf{v} = [\text{TVS}_1, \dots, \text{TVS}_{50}] \in \mathbb{R}^T$, one scalar per timestep, directly capturing how cross-seed semantic divergence evolves over the denoising schedule.

4.3 Classification via Logistic Regression

We classify the TVS feature vector using a **Logistic Regression** model (max_iter=2000, ℓ_2 regularization). This linear model achieves competitive accuracy, confirming the conflict signal is globally distributed across the dispersion curve.

5 Experimental Setup

5.1 Models and Datasets

We evaluate TVS on two architecturally distinct DLMs: **LLaDA-8B-Instruct** (Nie et al., 2025), which uses low-confidence remasking, and **Dream-v0-Instruct-7B** (Ye et al., 2025), which uses entropy-based remasking. This pairing allows us to assess whether TVS generalizes across different diffusion denoising strategies.

Our testbed comprises four datasets: **Synthetic** (custom queries generated via LLM templates to explicitly pair simple facts with direct contradictions), **SciQ** (Welbl et al., 2017), **PopQA** (Mallen et al., 2023), and **CounterFact** (Meng et al., 2022). We selected these because their factual format enables straightforward injection of contradictory information. We randomly select 3,000 instances per dataset. We filter these for instances where the base DLM correctly generates the true answer purely from parametric memory across $\geq 4/5$ seeds. This strict parametric filter heavily decimates datasets with lower baseline knowledge: of the 12,000 initial candidates, **2,947 facts** survived for LLaDA (SciQ: 1,448; CounterFact: 870; Synthetic: 364; PopQA: 265) and **2,370 facts** survived for Dream (reflecting Dream’s somewhat lower parametric recall). To simulate conflict, we inject a contradictory RAG context, yielding 5,894 (LLaDA) and 4,740 (Dream) balanced paired samples respectively. The final dataset for each model is pooled across all domains and uses a 70/15/15 stratified train/validation/test split.

5.2 Implementation Details

The evaluation pipeline is identical for both models. Trajectories were collected using $N = 2$ to 5 seeds over $T = 50$ timesteps. Sentence embeddings use the frozen all-mpnet-base-v2 model. The Logistic Regression is trained on the pooled, flattened $\mathbf{v} \in \mathbb{R}^{50}$ TVS feature vectors (max_iter=2000, ℓ_2 regularization). All results are reported as mean $\pm$ standard deviation over 5 independent random train/test splits.

6 Results

6.1 Overall Detection Performance

Table 1 presents the overall conflict detection performance of our TVS approach on both LLaDA and Dream 7B. As an informal point of reference, we also include a best-effort replication of TraceDet’s Variational Information Bottleneck trajectory classifier (the original codebase is not publicly available and has not been formally verified on this task). On LLaDA, our simple linear TVS classifier achieves $70.76\% \pm 0.62\%$ accuracy, performing similarly to this approximated baseline. On Dream 7B, TVS-LR achieves $63.91\% \pm 1.17\%$, confirming that the conflict signal transfers across architectures, albeit with reduced magnitude. Notably, TVS-LR slightly outperforms TraceDet on Dream

in AUROC (0.6896 vs. 0.6876), suggesting that the simpler embedding-based approach is comparably robust across architectures.

Method	LLaDA Acc (%)	LLaDA AUROC	Dream Acc (%)	Dream AUROC
TraceDet*	71.34 $\pm$ 0.66	0.7865 $\pm$ 0.0078	64.56 $\pm$ 0.56	0.6876 $\pm$ 0.0056
TVS-LR (Ours)	70.76 $\pm$ 0.62	0.7725 $\pm$ 0.0082	63.91 $\pm$ 1.17	0.6896 $\pm$ 0.0101

Table 1: Overall conflict detection performance (mean $\pm$ std over 5 random splits). TVS achieves competitive accuracy on both models. The $\sim 7\%$ gap on Dream likely reflects its smoother entropy-based remasking dynamics. *TraceDet codebase is not publicly available; we benchmark against an unverified best-effort replication for context.

Metric	LLaDA (%)	Dream (%)
Precision	72.95 $\pm$ 1.32	66.13 $\pm$ 0.91
Recall	66.05 $\pm$ 1.46	56.95 $\pm$ 3.10
F1-Score	69.30 $\pm$ 0.60	61.16 $\pm$ 2.00

Table 2: Classification metrics for Logistic Regression (mean $\pm$ std over 5 random splits). Dream’s lower recall (56.95%) indicates that its smoother trajectories make conflict samples harder to distinguish from clean ones.

6.2 Dataset-Specific Dynamics

Table 3 provides the per-domain accuracy breakdown (LR classifier, averaged over 5 random splits) for both models.

Model	Synth.	SciQ	PopQA	CounterFact
LLaDA	77.31 $\pm$ 5.92	74.52 $\pm$ 1.44	72.34 $\pm$ 1.35	61.80 $\pm$ 2.48
Dream	64.82 $\pm$ 4.24	65.63 $\pm$ 1.54	59.59 $\pm$ 3.94	58.16 $\pm$ 5.60

Table 3: Detection accuracy (%) by domain (LR classifier, mean $\pm$ std over 5 splits). The cross-dataset difficulty ordering (Synth. > SciQ > PopQA > CounterFact) is preserved across both architectures.

Both models perform best on Synthetic facts and worst on CounterFact, with the cross-dataset difficulty ordering fully preserved. The high performance on Synthetic is likely an artifact of its templated construction: structurally uniform conflicts produce a clean friction signal. The CounterFact difficulty likely arises because semantically extreme counterfactuals (e.g., “The capital of France is London”) override the parametric prior so decisively that cross-seed divergence collapses early, leaving a weak TVS signal.

Dream’s per-dataset accuracy is uniformly lower, with the largest gaps on Synthetic (-12.5%) and

PopQA (-12.8%). However, the consistent preservation of the difficulty ordering across two architecturally distinct DLMs—with fundamentally different remasking strategies—provides strong evidence that the relative strength of the TVS signal is driven by dataset properties (e.g., conflict type, phrasing complexity) rather than model-specific artifacts.

7 Ablation Studies

7.1 Architecture Comparison: Is Sequential Modeling Required?

Table 4 compares three classifier architectures over 5 random splits on the TVS feature vector for both models.

Classifier	LLaDA Acc (%)	LLaDA AUROC	Dream Acc (%)	Dream AUROC
Logistic Regression	71.05 $\pm$ 0.72	0.7701 $\pm$ 0.0098	63.40 $\pm$ 1.30	0.6972 $\pm$ 0.0072
LSTM (Uni-directional)	71.75 $\pm$ 0.54	0.7895 $\pm$ 0.0061	62.87 $\pm$ 1.14	0.6713 $\pm$ 0.0069
Bi-LSTM + Attention	71.93 $\pm$ 0.71	0.7906 $\pm$ 0.0062	64.67 $\pm$ 0.70	0.6977 $\pm$ 0.0069

Table 4: Architecture comparison over 5 random splits. On both models, all confidence intervals overlap substantially. The LR-to-Bi-LSTM gap is 0.88% on LLaDA and 1.27% on Dream, confirming that sequential modeling adds negligible benefit across architectures.

The marginal performance gap between the simple linear model and the complex recurrent architectures is consistent across both DLMs, suggesting that the classifier is primarily exploiting the overall magnitude of dispersion rather than complex temporal dependencies. To rigorously test this, we evaluated both the Logistic Regression and the LSTM on a dataset where the 50 timestep features were randomly shuffled (consistently across all samples). By construction, the Logistic Regression’s accuracy is invariant to feature permutation and remains identically unchanged. Empirically, the LSTM’s accuracy also did not drop under this shuffling constraint (remaining at $\sim 70.6\%$). This confirms that no sequential structure is being exploited by any model, and that recurrent architectures do not extract meaningful signal beyond what the order-invariant linear baseline already captures.

7.2 Seed Count Ablation: How Many Runs Are Needed?

Table 5 evaluates the effect of reducing the number of parallel denoising runs N , using the Logistic Regression classifier.

The two models exhibit strikingly different seed sensitivities. LLaDA’s accuracy is essentially flat from $N = 2$ to $N = 5$ ($+0.47\%$, within error bars), confirming that two parallel runs are sufficient. In

N Seeds	LLaDA Acc (%)	LLaDA AUROC	Dream Acc (%)	Dream AUROC
2	70.76 $\pm$ 0.62	0.7725 $\pm$ 0.0082	63.91 $\pm$ 1.17	0.6893 $\pm$ 0.0119
3	71.75 $\pm$ 1.06	0.7806 $\pm$ 0.0084	67.12 $\pm$ 0.95	0.7261 $\pm$ 0.0131
5	71.23 $\pm$ 1.00	0.7851 $\pm$ 0.0074	69.62 $\pm$ 0.51	0.7589 $\pm$ 0.0134

Table 5: Effect of seed count on detection performance (LR classifier, 5 random splits). LLaDA’s accuracy is stable across seed counts ($\Delta = +0.47\%$), while Dream benefits substantially from additional seeds ($\Delta = +5.71\%$), suggesting the optimal seed count is architecture-dependent.

contrast, Dream benefits substantially: accuracy rises from 63.91% to 69.62% (+5.71%), nearly closing the cross-architecture gap. This asymmetry is likely attributable to Dream’s entropy-based remasking, which produces smoother denoising trajectories with less per-run variance. With only $N = 2$ seeds, the pairwise cosine distance is computed from a single pair, making the signal highly sensitive to stochastic noise. Increasing to $N = 5$ yields $\binom{5}{2} = 10$ pairwise comparisons, substantially improving the robustness of the TVS estimate and amplifying the conflict signal in Dream’s smoother trajectory space.

8 Analysis: Interpreting the Tug-of-War

Figure 2 plots the average Cosine Dispersion (TVS) over the 50-step denoising trajectory for LLaDA, revealing that knowledge friction manifests as a sustained, three-phase elevation in cross-seed semantic divergence.




Figure 2: Average TVS across the 50-step denoising trajectory for LLaDA (Clean: blue solid; RAG Conflict: red dashed; shaded = SEM). Three phases are visible: **Zone 1** ($t \in [0, 15]$): Early Exploration—both conditions exhibit high dispersion. **Zone 2** ($t \in [15, 28]$): Mid-Step Friction—the clean curve decays while the conflict curve remains elevated (gap ≈ 0.15). **Zone 3** ($t \in [28, 50]$): Delayed Crystallization—both curves collapse, but a residual gap persists ($\Delta \approx 0.07$).

On LLaDA, the conflict trajectory (red dashed) starts at CD ≈ 0.51 versus 0.39 for clean (blue




Figure 3: Average TVS across the 50-step denoising trajectory for Dream 7B ($N = 2$ seeds; Clean: blue solid; RAG Conflict: red dashed; shaded = SEM). The same three-phase structure is observed, but with a narrower clean-vs-conflict gap, consistent with Dream’s smoother entropy-based remasking dynamics.




Figure 4: Detailed TVS dynamics for LLaDA with $\pm$ std shading ($N = 2$ seeds). The conflict curve (red) starts at ≈ 0.51 vs. clean (green) at ≈ 0.39 , maintaining a consistent gap through $t \approx 28$ before both converge—with conflict retaining a higher residual dispersion (≈ 0.07 vs. ≈ 0.01).

solid), maintaining a gap of ~ 0.1 – 0.15 throughout Zone 1 and Zone 2. Both curves collapse rapidly after $t \approx 28$, but a residual gap of ≈ 0.07 persists at convergence, representing the “crystallization gap” that distinguishes conflict samples even after generation terminates.

Critically, Dream 7B (Figure 3) exhibits the same three-phase temporal structure, providing qualitative cross-architecture validation of the TVS framework. The clean-vs-conflict separation is narrower on Dream, consistent with its entropy-based remasking producing smoother, less oscillatory trajectories. This narrower gap directly explains Dream’s lower classification accuracy (Table 1) and its greater sensitivity to seed count (Table 5): the weaker per-pair signal requires more parallel runs to surface reliably.

This three-phase structure directly motivates the

temporal feature representation: a simple classifier over the dispersion curve captures the sustained divergence without requiring a sequential model.

8.1 The Value of the Temporal Trajectory

To explicitly address whether the intermediate diffusion trajectory provides signal beyond standard output-level multi-sample variance (analogous to autoregressive semantic entropy), we evaluated a Logistic Regression trained strictly on the final crystallization step (TVS₅₀). This trivial baseline achieved only $60.14\% \pm 0.57\%$ accuracy. The $\sim 11\%$ performance drop compared to the full T -dimensional TVS vector (71.05%) empirically proves that the mid-step “tug-of-war” dynamics contain critical diagnostic signal for detecting knowledge friction that is lost if one only evaluates the final generated output.

9 Conclusion

We introduced the Trajectory Variance Score (TVS), a simple and interpretable method for detecting RAG-induced knowledge conflicts in Discrete Diffusion Language Models. By computing mean pairwise cosine distance across $N = 2$ stochastic denoising trajectories at each timestep, TVS produces a T -dimensional feature vector. We demonstrate that a simple Logistic Regression classifier over this vector achieves competitive detection accuracy, proving that no complex sequential modeling is required to extract the conflict signal. The TVS curve directly visualizes the temporal tug-of-war between parametric and contextual attractors as a three-phase dynamics profile. Evaluation across two architecturally distinct DLMs—LLaDA (low-confidence remasking) and Dream 7B (entropy-based remasking)—demonstrates that the TVS signal, the three-phase temporal dynamics, and the sufficiency of linear classifiers all generalize across diffusion architectures. The architecture-dependent seed sensitivity we observe (Dream benefits substantially from $N > 2$, while LLaDA does not) provides a practical guideline: the optimal seed count should be calibrated per model. With as few as two parallel inference runs, TVS is a computationally lightweight and interpretable diagnostic for investigating knowledge friction in diffusion language models.

10 Limitations

TVS requires multiple parallel denoising runs, which introduces inference overhead compared to single-pass methods. Additionally, our pipeline relies on exact sub-string matching to isolate answer entities; this rigid extraction step is a potential failure point for valid paraphrases or partial matches, leading to unquantified extraction failure rates. A further limitation is that TVS relies on a parametric memory filter that itself requires $N = 5$ generation runs at evaluation time; exploring lighter filtering strategies is a promising direction for future work. Finally, while we demonstrate generality across two DLMs, TVS accuracy is architecture-dependent: Dream 7B achieves $\sim 64\%$ accuracy with $N = 2$ seeds versus LLaDA’s $\sim 71\%$, and models with smoother denoising dynamics may require more parallel runs to surface a reliable conflict signal. Whether TVS generalizes to DLMs with substantially different architectures or training regimes remains an open question.

Ethics Statement

As language models are increasingly deployed in real-world applications via Retrieval-Augmented Generation, their susceptibility to knowledge conflicts poses significant risks of propagating misinformation or adversarial fakes. The Trajectory Variance Score (TVS) provides an interpretability mechanism to detect and intercept these hallucinations before they reach the end user. While TVS significantly enhances the transparency of discrete diffusion models, it should not be treated as a standalone safeguard. It is intended to complement, rather than replace, standard fact-checking pipelines, ensuring that AI systems remain trustworthy and aligned with factual reality.

References

Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. 2023. [Discovering latent knowledge in language models without supervision](#). In *International Conference on Learning Representations (ICLR)*.

Shenxu Chang, Junchi Yu, Weixing Wang, Yongqiang Chen, Jialin Yu, Philip Torr, and Jindong Gu. 2025. Tracedet: Hallucination detection in discrete diffusion models via hidden trajectories. *arXiv preprint arXiv:2510.01274*.

Arshia Hemmat, Philip Torr, Yongqiang Chen, and Junchi Yu. 2026. Tdgnnet: Hallucination detection

535	in diffusion language models via temporal dynamic	Chuanyue Yu, Jiahui Wang, Yuhan Li, Heng Chang,	587
536	graphs. <i>arXiv preprint arXiv:2602.08048</i> .	Ge Lan, Qingyun Sun, Jia Li, Jianxin Li, and Ziwei	588
		Zhang. 2026. Unlocking the potentials of retrieval-	589
537	Jaemin Kim and Jong Chul Ye. 2026. Adaptive guid-	augmented generation for diffusion language models.	590
538	ance for retrieval-augmented masked diffusion mod-	<i>arXiv preprint arXiv:2601.11342</i> .	591
539	els. <i>arXiv preprint arXiv:2603.17677</i> .		
540	Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. 2023.		
541	Semantic uncertainty: Linguistic invariances for un-		
542	certainty estimation in natural language generation .		
543	In <i>The Eleventh International Conference on Learn-</i>		
544	<i>ing Representations</i> .		
545	Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio		
546	Petroni, Vladimir Karpukhin, Naman Goyal, Hein-		
547	rich Küttler, Mike Lewis, Wen-tau Yih, Tim Rock-		
548	täschel, Sebastian Riedel, and Douwe Kiela. 2020.		
549	Retrieval-augmented generation for knowledge-		
550	intensive nlp tasks. <i>Advances in Neural Information</i>		
551	<i>Processing Systems</i> .		
552	Alex Mallen, Akari Asai, Victor Zhong, Rajarshi Das,		
553	Daniel Khashabi, and Hannaneh Hajishirzi. 2023.		
554	When not to trust language models: Investigating		
555	effectiveness of parametric and non-parametric mem-		
556	ories . In <i>Proceedings of the 61st Annual Meeting of</i>		
557	<i>the Association for Computational Linguistics (ACL)</i> .		
558	Kevin Meng, David Bau, Alex Andonian, and Yonatan		
559	Belinkov. 2022. Locating and editing factual asso-		
560	ciations in GPT . <i>Advances in Neural Information</i>		
561	<i>Processing Systems</i> .		
562	Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang,		
563	Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-Rong		
564	Wen, and Chongxuan Li. 2025. Large language dif-		
565	fusion models. <i>arXiv preprint arXiv:2502.09992</i> .		
566	Subham Sekhar Sahoo, M. Arriola, Y. Schiff,		
567	A. Gokaslan, E. Marroquin, J. Chiu, A. Rush, and		
568	V. Kuleshov. 2024. Simple and effective masked		
569	diffusion language models . In <i>Advances in Neural</i>		
570	<i>Information Processing Systems</i> .		
571	Yash Shah, Abhijit Chakraborty, Naresh Kumar Devu-		
572	lapally, Vishnu Lokhande, and Vivek Gupta. 2026.		
573	Oscar: Orchestrated self-verification and cross-path		
574	refinement. <i>arXiv preprint arXiv:2604.01624</i> .		
575	Jiaxin Shi, Kehang Han, Zhe Wang, Arnaud Doucet, and		
576	Michalis K. Titsias. 2024. Simplified and generalized		
577	masked diffusion for discrete data . In <i>Advances in</i>		
578	<i>Neural Information Processing Systems</i> .		
579	Johannes Welbl, Nelson F Liu, and Matt Gardner. 2017.		
580	Crowdsourcing multiple choice science questions .		
581	In <i>Proceedings of the 3rd Workshop on Noisy User-</i>		
582	<i>generated Text</i> .		
583	Jiacheng Ye, Zhihui Xie, Lin Zheng, Jiahui Gao, Zirui		
584	Wu, Xin Jiang, Zhenguo Li, and Lingpeng Kong.		
585	2025. Dream 7b: Diffusion large language models.		
586	<i>arXiv preprint arXiv:2508.15487</i> .		